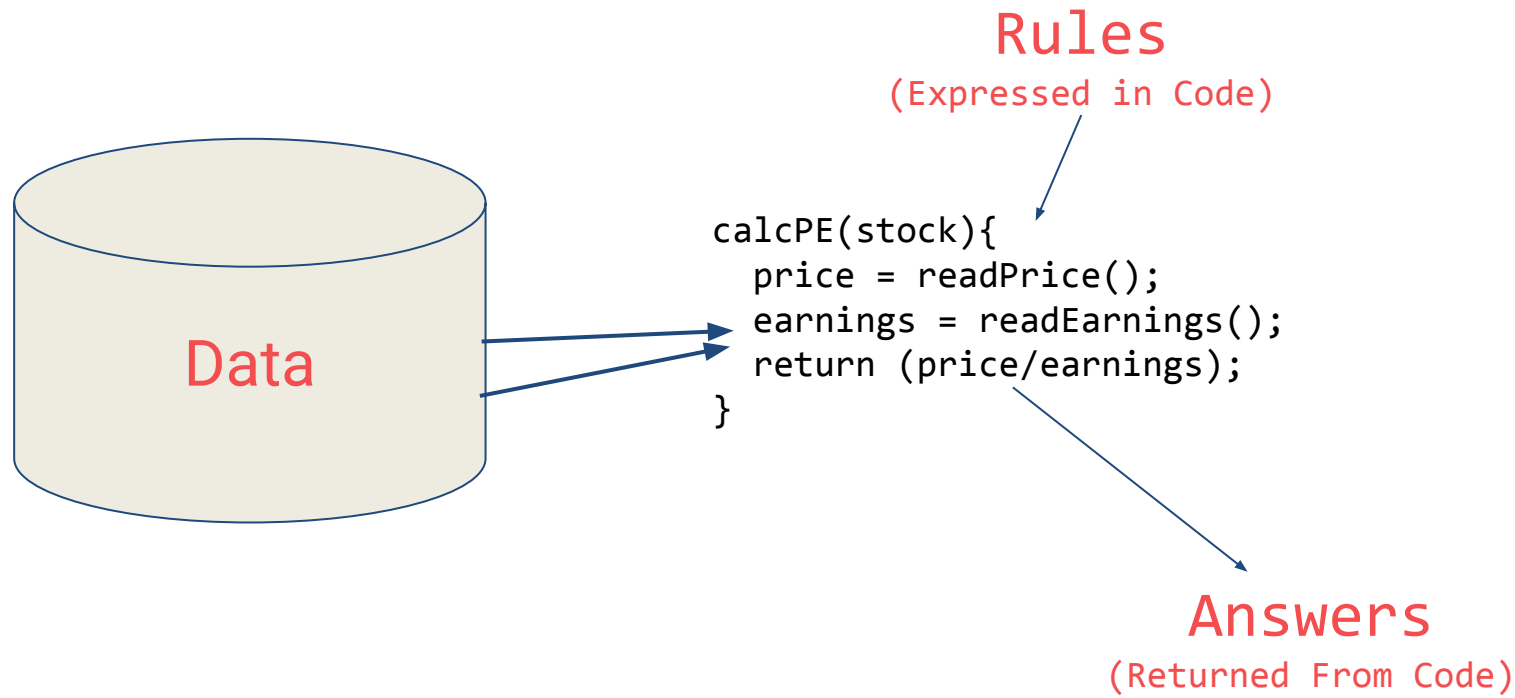


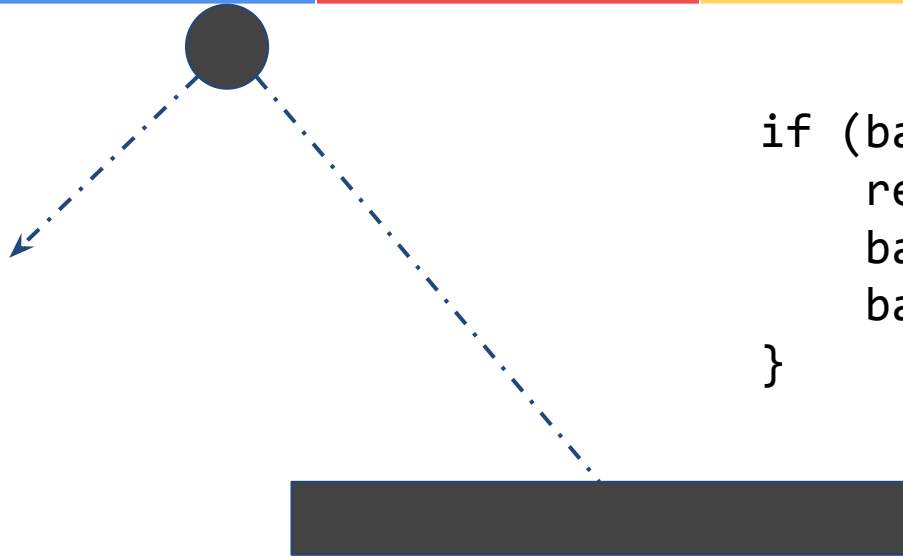
Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see <https://creativecommons.org/licenses/by-sa/2.0/legalcode>





```
if (ball.collide(brick)){  
    removeBrick();  
    ball.dx=-1*(ball.dx);  
    ball.dy=-1*(ball.dy);  
}
```







Activity Recognition



```
if(speed<4){  
    status=WALKING;  
}
```

Activity Recognition



```
if(speed<4){  
    status=WALKING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else {  
    status=RUNNING;  
}
```



Activity Recognition



```
if(speed<4){  
    status=WALKING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else {  
    status=RUNNING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else if(speed<12){  
    status=RUNNING;  
} else {  
    status=BIKING;  
}
```


Activity Recognition



```
if(speed<4){  
    status=WALKING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else {  
    status=RUNNING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else if(speed<12){  
    status=RUNNING;  
} else {  
    status=BIKING;  
}
```



// Oh crap



Activity Recognition



0101001010100101010
1001010101001011101
0100101010010101001
0101001010100101010

Label = WALKING



1010100101001010101
0101010010010010001
0010011111010101111
1010100100111101011

Label = RUNNING



10010100111111010101
1101010111010101110
1010101111010101011
1111110001111010101

Label = BIKING



11111111111010011101
0011111010111110101
0101110101010101110
1010101010100111110

Label = GOLFING
(Sort of)

$X = -1, 0, 1, 2, 3, 4$

$Y = -3, -1, 1, 3, 5, 7$



```
model = tf.keras.Sequential([  
    tf.keras.Input(shape=(1,)),  
    tf.keras.layers.Dense(units=1)  
])
```



```
model = tf.keras.Sequential([  
    tf.keras.Input(shape=(1,)),  
    tf.keras.layers.Dense(units=1)  
])
```



```
model = tf.keras.Sequential([  
    tf.keras.Input(shape=(1,)),  
    tf.keras.layers.Dense(units=1)  
])
```



```
model = tf.keras.Sequential([  
    tf.keras.Input(shape=(1,)),  
    tf.keras.layers.Dense(units=1)  
])
```




```
model = tf.keras.Sequential([  
    tf.keras.Input(shape=(1,)),  
    tf.keras.layers.Dense(units=1)  
])  
model.compile(optimizer='sgd', loss='mean_squared_error')
```

```
model = tf.keras.Sequential([
    tf.keras.Input(shape=(1,)),
    tf.keras.layers.Dense(units=1)
])
model.compile(optimizer='sgd', loss='mean_squared_error')

xs = np.array([-1.0, 0.0, 1.0, 2.0, 3.0, 4.0], dtype=float)
ys = np.array([-3.0, -1.0, 1.0, 3.0, 5.0, 7.0], dtype=float)
```



```
model = tf.keras.Sequential([
    tf.keras.Input(shape=(1,)),
    tf.keras.layers.Dense(units=1)
])
model.compile(optimizer='sgd', loss='mean_squared_error')

xs = np.array([-1.0, 0.0, 1.0, 2.0, 3.0, 4.0], dtype=float)
ys = np.array([-3.0, -1.0, 1.0, 3.0, 5.0, 7.0], dtype=float)

model.fit(xs, ys, epochs=500)
```



```
model = tf.keras.Sequential([
    tf.keras.Input(shape=(1,)),
    tf.keras.layers.Dense(units=1)
])
model.compile(optimizer='sgd', loss='mean_squared_error')

xs = np.array([-1.0, 0.0, 1.0, 2.0, 3.0, 4.0], dtype=float)
ys = np.array([-3.0, -1.0, 1.0, 3.0, 5.0, 7.0], dtype=float)

model.fit(xs, ys, epochs=500)

model.predict(np.array([10.0]))
```



Copyright Notice

These slides are distributed under the Creative Commons License.


[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see <https://creativecommons.org/licenses/by-sa/2.0/legalcode>




Fashion MNIST

- 70k Images
- 10 Categories
- Images are 28x28
- Can train a neural net!




Fashion MNIST

- 70k Images
- 10 Categories
- Images are 28x28
- Can train a neural net!





```
fashion_mnist = tf.keras.datasets.fashion_mnist  
(train_images, train_labels), (test_images, test_labels) = fashion_mnist.load_data()
```

```
fashion_mnist = tf.keras.datasets.fashion_mnist  
(train_images, train_labels), (test_images, test_labels) = fashion_mnist.load_data()
```



09

```
fashion_mnist = tf.keras.datasets.fashion_mnist  
(train_images, train_labels), (test_images, test_labels) = fashion_mnist.load_data()
```



09

09 = ankle boot;
踝靴;
アンクルブーツ;
Bróg rúitín

```
model = tf.keras.Sequential([  
    tf.keras.Input(shape=(28, 28)),  
    tf.keras.layers.Flatten(),  
    tf.keras.layers.Dense(128, activation=tf.nn.relu),  
    tf.keras.layers.Dense(10, activation=tf.nn.softmax)  
])
```



```
model = tf.keras.Sequential([
    tf.keras.Input(shape=(28, 28)),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation=tf.nn.relu),
    tf.keras.layers.Dense(10, activation=tf.nn.softmax)
])
```





```
model = tf.keras.Sequential([
    tf.keras.Input(shape=(28, 28)),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation=tf.nn.relu),
    tf.keras.layers.Dense(10, activation=tf.nn.softmax)
])
```



```
model = tf.keras.Sequential([  
    tf.keras.Input(shape=(28, 28)),  
    tf.keras.layers.Flatten(),  
    tf.keras.layers.Dense(128, activation=tf.nn.relu),  
    tf.keras.layers.Dense(10, activation=tf.nn.softmax)  
])
```




Input




$$w_0x_0 + w_1x_1 + w_2x_2 \dots w_Nx_N = 9$$

Output



What is a Neural Network?









Neural Network Overview (C1W3L01)

11,067 views


43 0 SHARE SAVE ...

neural Networks and Deep Learning (Course 1 of
Deeplearning.ai - 25 / 43

≡

- ▶  **Neural Network Overview (C1W3L01)**
Deeplearning.ai
4:27
- 26  **Neural Network Representations (C1W3L02)**
Deeplearning.ai
5:15
- 27  **Computing Neural Network Output (C1W3L03)**
Deeplearning.ai
9:58
- 28  **Vectorizing Across Multiple Examples (C1W3L04)**
Deeplearning.ai
9:06
- 29  **Explanation For Vectorized Implementation (C1W3L05)**
Deeplearning.ai
7:38
- 30  **Activation Functions (C1W3L06)**
Deeplearning.ai
10:57

Why Non-linear Activation Functions

 Login
Username
Password
Sign Up
Not yet a user? Sign up

Complete User Registration system using PHP and MySQL...
Awa Melvine
5.7M views
32:43

```
fashion_mnist = tf.keras.datasets.fashion_mnist
(train_images, train_labels), (test_images, test_labels) = fashion_mnist.load_data()
training_images = training_images / 255.0
test_images = test_images / 255.0
model = tf.keras.Sequential([
    tf.keras.Input(shape=(28, 28)),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation=tf.nn.relu),
    tf.keras.layers.Dense(10, activation=tf.nn.softmax)
])
model.compile(optimizer=tf.optimizers.Adam(), loss='sparse_categorical_crossentropy')
model.fit(training_images, training_labels, epochs=5)
```



```
fashion_mnist = tf.keras.datasets.fashion_mnist
(train_images, train_labels), (test_images, test_labels) = fashion_mnist.load_data()
training_images = training_images / 255.0
test_images = test_images / 255.0
model = tf.keras.Sequential([
    tf.keras.Input(shape=(28, 28)),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation=tf.nn.relu),
    tf.keras.layers.Dense(10, activation=tf.nn.softmax)
])
model.compile(optimizer=tf.optimizers.Adam(), loss='sparse_categorical_crossentropy')
model.fit(training_images, training_labels, epochs=5)
```



```
class myCallback(tf.keras.callbacks.Callback):  
    def on_epoch_end(self, epoch, logs=None):  
        if logs['loss'] < 0.4:  
            print('Loss is low so cancelling training!')  
            self.model.stop_training = True
```



```
class myCallback(tf.keras.callbacks.Callback):  
    def on_epoch_end(self, epoch, logs=None):  
        if logs['loss'] < 0.4:  
            print('Loss is low so cancelling training!')  
            self.model.stop_training = True
```



```
class myCallback(tf.keras.callbacks.Callback):  
    def on_epoch_end(self, epoch, logs=None):  
        if logs['loss'] < 0.4:  
            print('Loss is low so cancelling training!')  
            self.model.stop_training = True
```



```
fashion_mnist = tf.keras.datasets.fashion_mnist
(train_images, train_labels), (test_images, test_labels) = fashion_mnist.load_data()
training_images = training_images / 255.0
test_images = test_images / 255.0
model = tf.keras.Sequential([
    tf.keras.input(shape=(28, 28)),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation=tf.nn.relu),
    tf.keras.layers.Dense(10, activation=tf.nn.softmax)
])
model.compile(optimizer=tf.optimizers.Adam(), loss='sparse_categorical_crossentropy')
model.fit(training_images, training_labels, epochs=5)
```



```
fashion_mnist = tf.keras.datasets.fashion_mnist
(train_images, train_labels), (test_images, test_labels) = fashion_mnist.load_data()
training_images = training_images / 255.0
test_images = test_images / 255.0
model = tf.keras.Sequential([
    tf.keras.input(shape=(28, 28)),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation=tf.nn.relu),
    tf.keras.layers.Dense(10, activation=tf.nn.softmax)
])
model.compile(optimizer=tf.optimizers.Adam(), loss='sparse_categorical_crossentropy')
model.fit(training_images, training_labels, epochs=5, callbacks=[myCallback()])
```

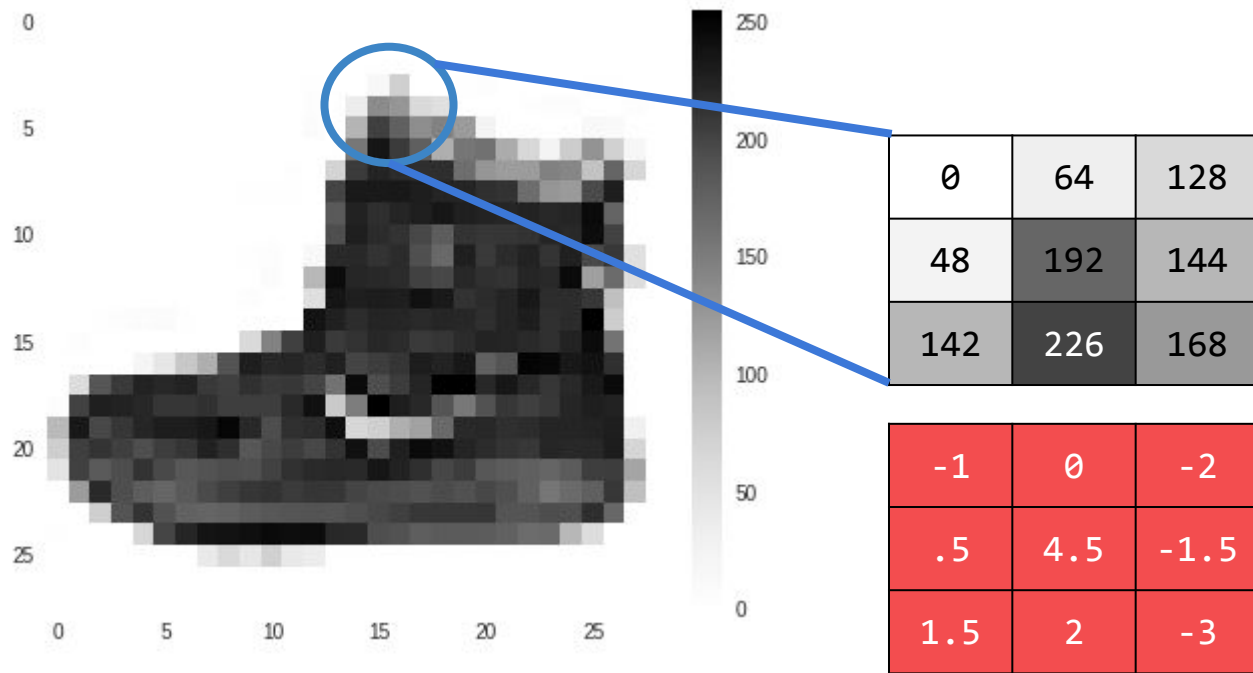


Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see <https://creativecommons.org/licenses/by-sa/2.0/legalcode>



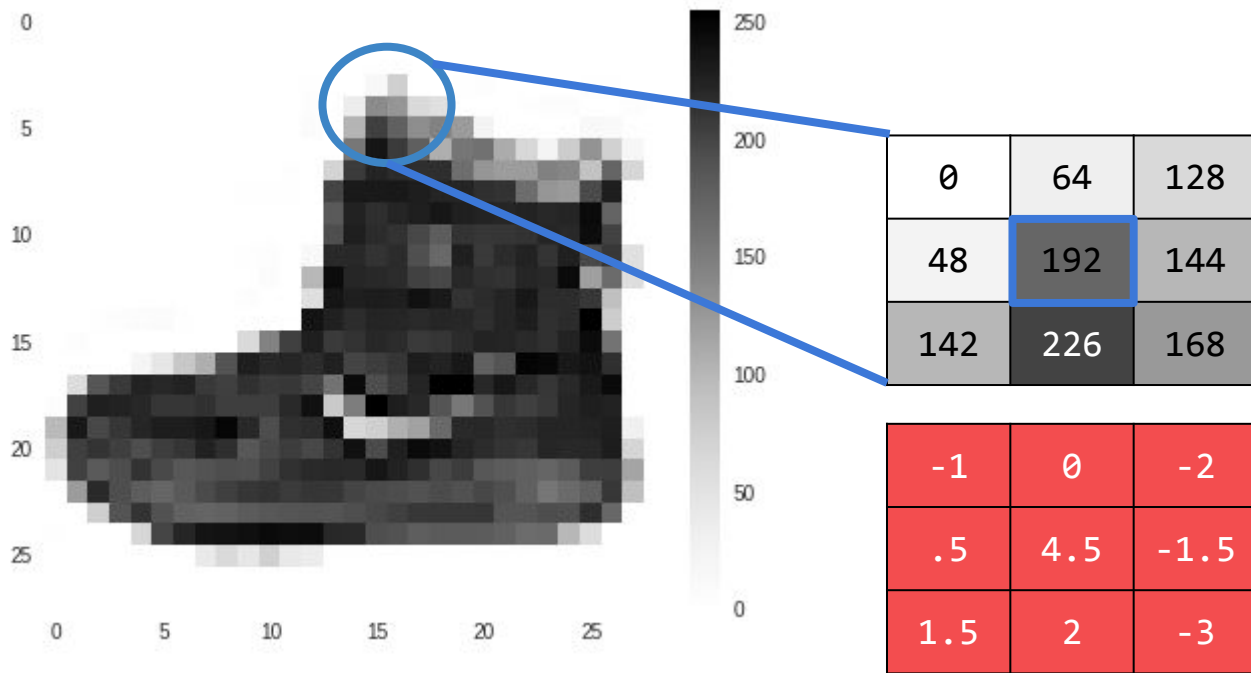
Current Pixel Value is 192

Consider neighbor Values

Filter Definition

CURRENT_PIXEL_VALUE = 192

NEW_PIXEL_VALUE = $(-1 * 0) + (0 * 64) + (-2 * 128) +$
 $(.5 * 48) + (4.5 * 192) + (-1.5 * 144) +$
 $(1.5 * 142) + (2 * 226) + (-3 * 168)$



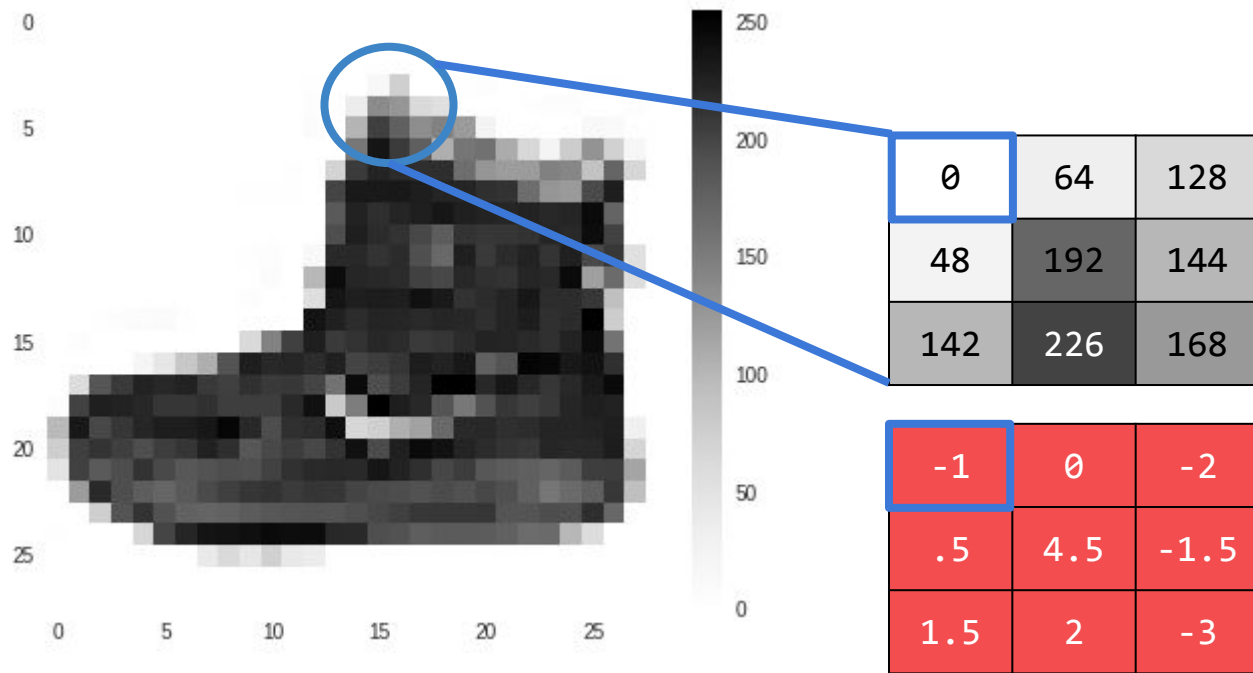
Current Pixel Value is 192

Consider neighbor Values

Filter Definition

CURRENT_PIXEL_VALUE = 192

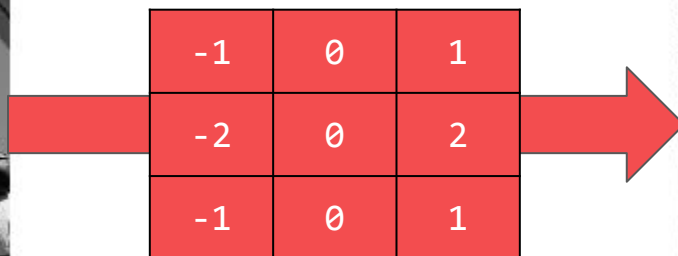
NEW_PIXEL_VALUE = $(-1 * 0) + (0 * 64) + (-2 * 128) +$
 $(.5 * 48) + (4.5 * 192) + (-1.5 * 144) +$
 $(1.5 * 142) + (2 * 226) + (-3 * 168)$

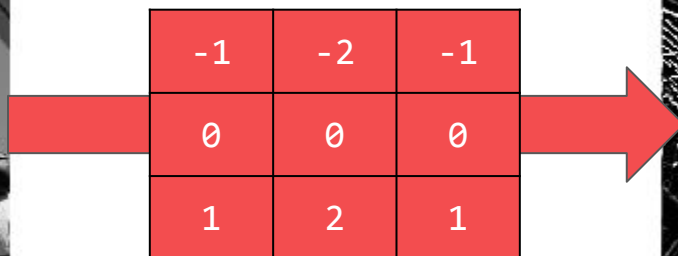


CURRENT_PIXEL_VALUE = 192

$$\begin{aligned} \text{NEW_PIXEL_VALUE} = & (-1 * 0) + (0 * 64) + (-2 * 128) + \\ & (.5 * 48) + (4.5 * 192) + (-1.5 * 144) + \\ & (1.5 * 142) + (2 * 226) + (-3 * 168) \end{aligned}$$







0	64	128	128
48	192	144	144
142	226	168	0
255	0	0	64

0	64
48	192

192

128	128
144	144

144

142	226
255	0

255

168	0
0	64

168

192	144
255	168




```
model = tf.keras.Sequential([
    tf.keras.Input(shape=(28, 28)),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation=tf.nn.relu),
    tf.keras.layers.Dense(10, activation=tf.nn.softmax)
])
```




```
model = tf.keras.Sequential([
    tf.keras.Input(shape=(28, 28, 1)),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')
])
```

```
model = tf.keras.Sequential([
    tf.keras.Input(shape=(28, 28, 1)),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')
])
```

Computer Vision Problems



Convolutional Neural Networks (Course 4 of the Deep Learning Specialization)

42 videos • 415,722 views • Last updated on Nov 7, 2017



Deeplearning.ai

SUBSCRIBE 28K


1



C4W1L01 Computer Vision

Deeplearning.ai


2



C4W1L02 Edge Detection Examples

Deeplearning.ai


3



C4W1L03 More Edge Detection

Deeplearning.ai


4



C4W1L04 Padding

Deeplearning.ai


5



C4W1L05 Strided Convolutions

Deeplearning.ai

6



C4W1L06 Convolutions Over Volumes

Deeplearning.ai

<https://bit.ly/2UGa7uH>



deeplearning.ai

```
model = tf.keras.Sequential([
    tf.keras.Input(shape=(28, 28, 1)),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')
])
```

```
model = tf.keras.Sequential([
    tf.keras.Input(shape=(28, 28, 1)),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')
])
```

```
model.summary()
```



Layer (type)	Output Shape	Param #
conv2d_12 (Conv2D)	(None, 26, 26, 64)	640
max_pooling2d_12 (MaxPooling)	(None, 13, 13, 64)	0
conv2d_13 (Conv2D)	(None, 11, 11, 64)	36928
max_pooling2d_13 (MaxPooling)	(None, 5, 5, 64)	0
flatten_5 (Flatten)	(None, 1600)	0
dense_10 (Dense)	(None, 128)	204928
dense_11 (Dense)	(None, 10)	1290

Layer (type)	Output Shape	Param #
conv2d_12 (Conv2D)	(None, 26, 26, 64)	640
max_pooling2d_12 (MaxPooling)	(None, 13, 13, 64)	0
conv2d_13 (Conv2D)	(None, 11, 11, 64)	36928
max_pooling2d_13 (MaxPooling)	(None, 5, 5, 64)	0
flatten_5 (Flatten)	(None, 1600)	0
dense_10 (Dense)	(None, 128)	204928
dense_11 (Dense)	(None, 10)	1290











Layer (type)	Output Shape	Param #
conv2d_12 (Conv2D)	(None, 26, 26, 64)	640
max_pooling2d_12 (MaxPooling)	(None, 13, 13, 64)	0
conv2d_13 (Conv2D)	(None, 11, 11, 64)	36928
max_pooling2d_13 (MaxPooling)	(None, 5, 5, 64)	0
flatten_5 (Flatten)	(None, 1600)	0
dense_10 (Dense)	(None, 128)	204928
dense_11 (Dense)	(None, 10)	1290

Layer (type)	Output Shape	Param #
conv2d_12 (Conv2D)	(None, 26, 26, 64)	640
max_pooling2d_12 (MaxPooling)	(None, 13, 13, 64)	0
conv2d_13 (Conv2D)	(None, 11, 11, 64)	36928
max_pooling2d_13 (MaxPooling)	(None, 5, 5, 64)	0
flatten_5 (Flatten)	(None, 1600)	0
dense_10 (Dense)	(None, 128)	204928
dense_11 (Dense)	(None, 10)	1290

Layer (type)	Output Shape	Param #
conv2d_12 (Conv2D)	(None, 26, 26, 64)	640
max_pooling2d_12 (MaxPooling)	(None, 13, 13, 64)	0
conv2d_13 (Conv2D)	(None, 11, 11, 64)	36928
max_pooling2d_13 (MaxPooling)	(None, 5, 5, 64)	0
flatten_5 (Flatten)	(None, 1600)	0
dense_10 (Dense)	(None, 128)	204928
dense_11 (Dense)	(None, 10)	1290

Layer (type)	Output Shape	Param #
conv2d_12 (Conv2D)	(None, 26, 26, 64)	640
max_pooling2d_12 (MaxPooling)	(None, 13, 13, 64)	0
conv2d_13 (Conv2D)	(None, 11, 11, 64)	36928
max_pooling2d_13 (MaxPooling)	(None, 5, 5, 64)	0
flatten_5 (Flatten)	(None, 1600)	0
dense_10 (Dense)	(None, 128)	204928
dense_11 (Dense)	(None, 10)	1290


Layer (type)	Output Shape	Param #
conv2d_12 (Conv2D)	(None, 26, 26, 64)	640
max_pooling2d_12 (MaxPooling)	(None, 13, 13, 64)	0
conv2d_13 (Conv2D)	(None, 11, 11, 64)	36928
max_pooling2d_13 (MaxPooling)	(None, 5, 5, 64)	0
flatten_5 (Flatten)	(None, 1600)	0
dense_10 (Dense)	(None, 128)	204928
dense_11 (Dense)	(None, 10)	1290

Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see <https://creativecommons.org/licenses/by-sa/2.0/legalcode>

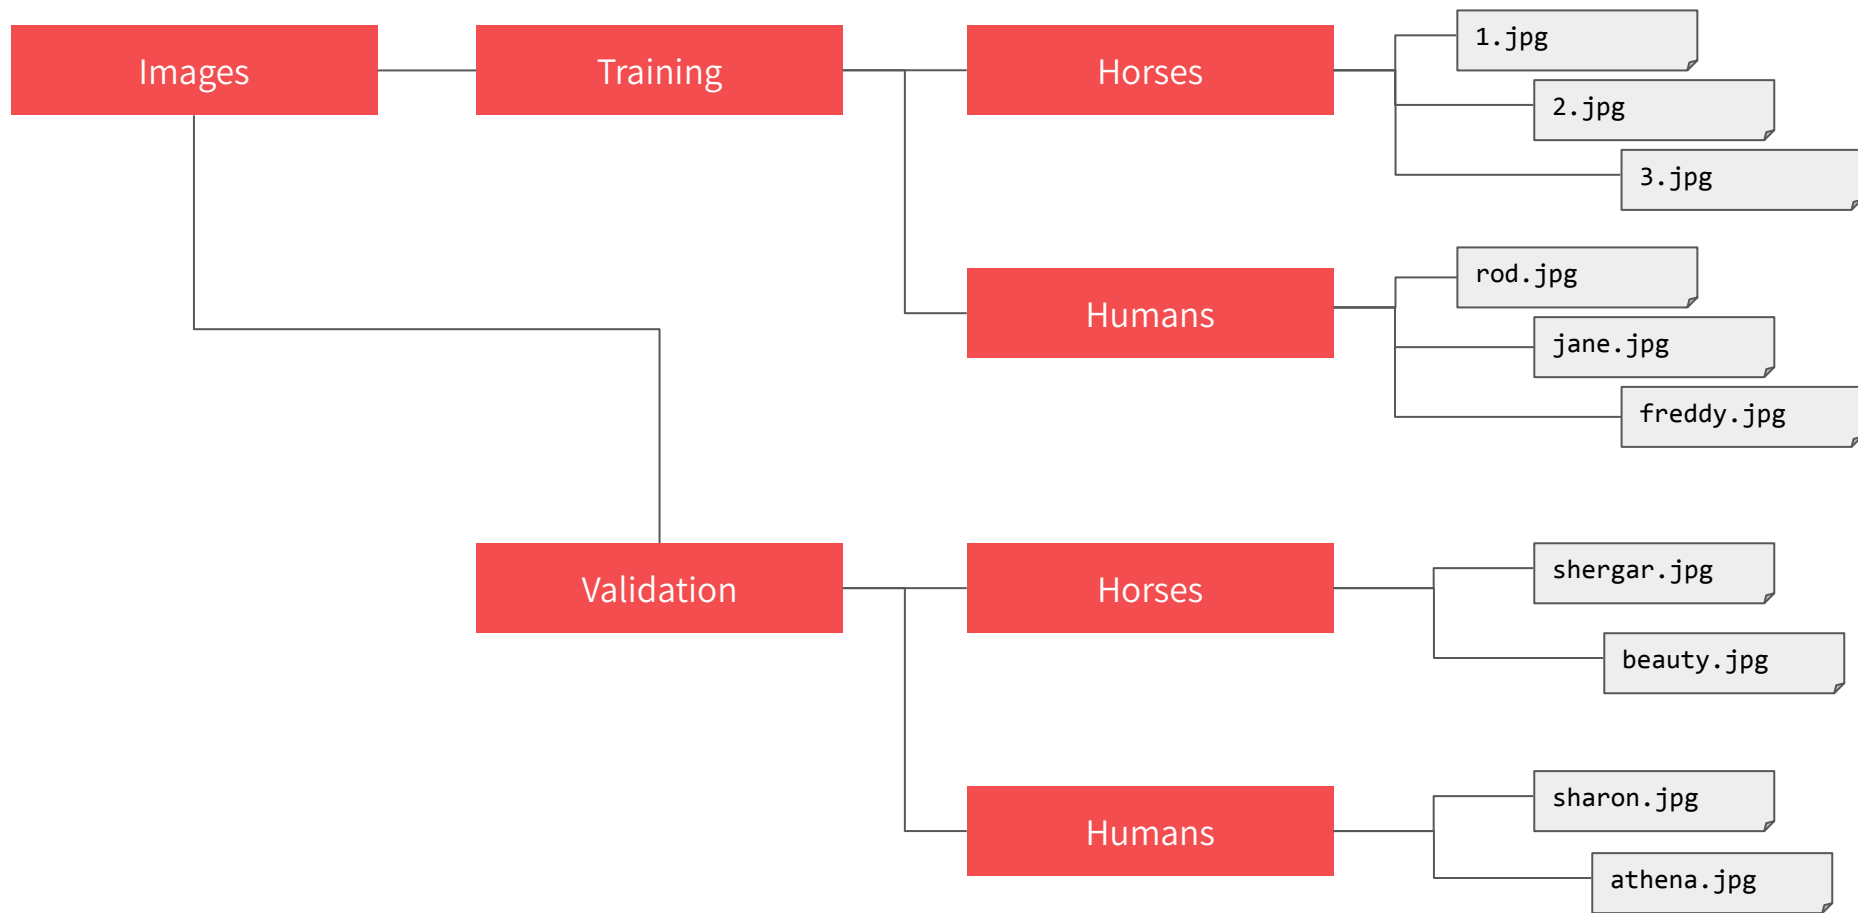


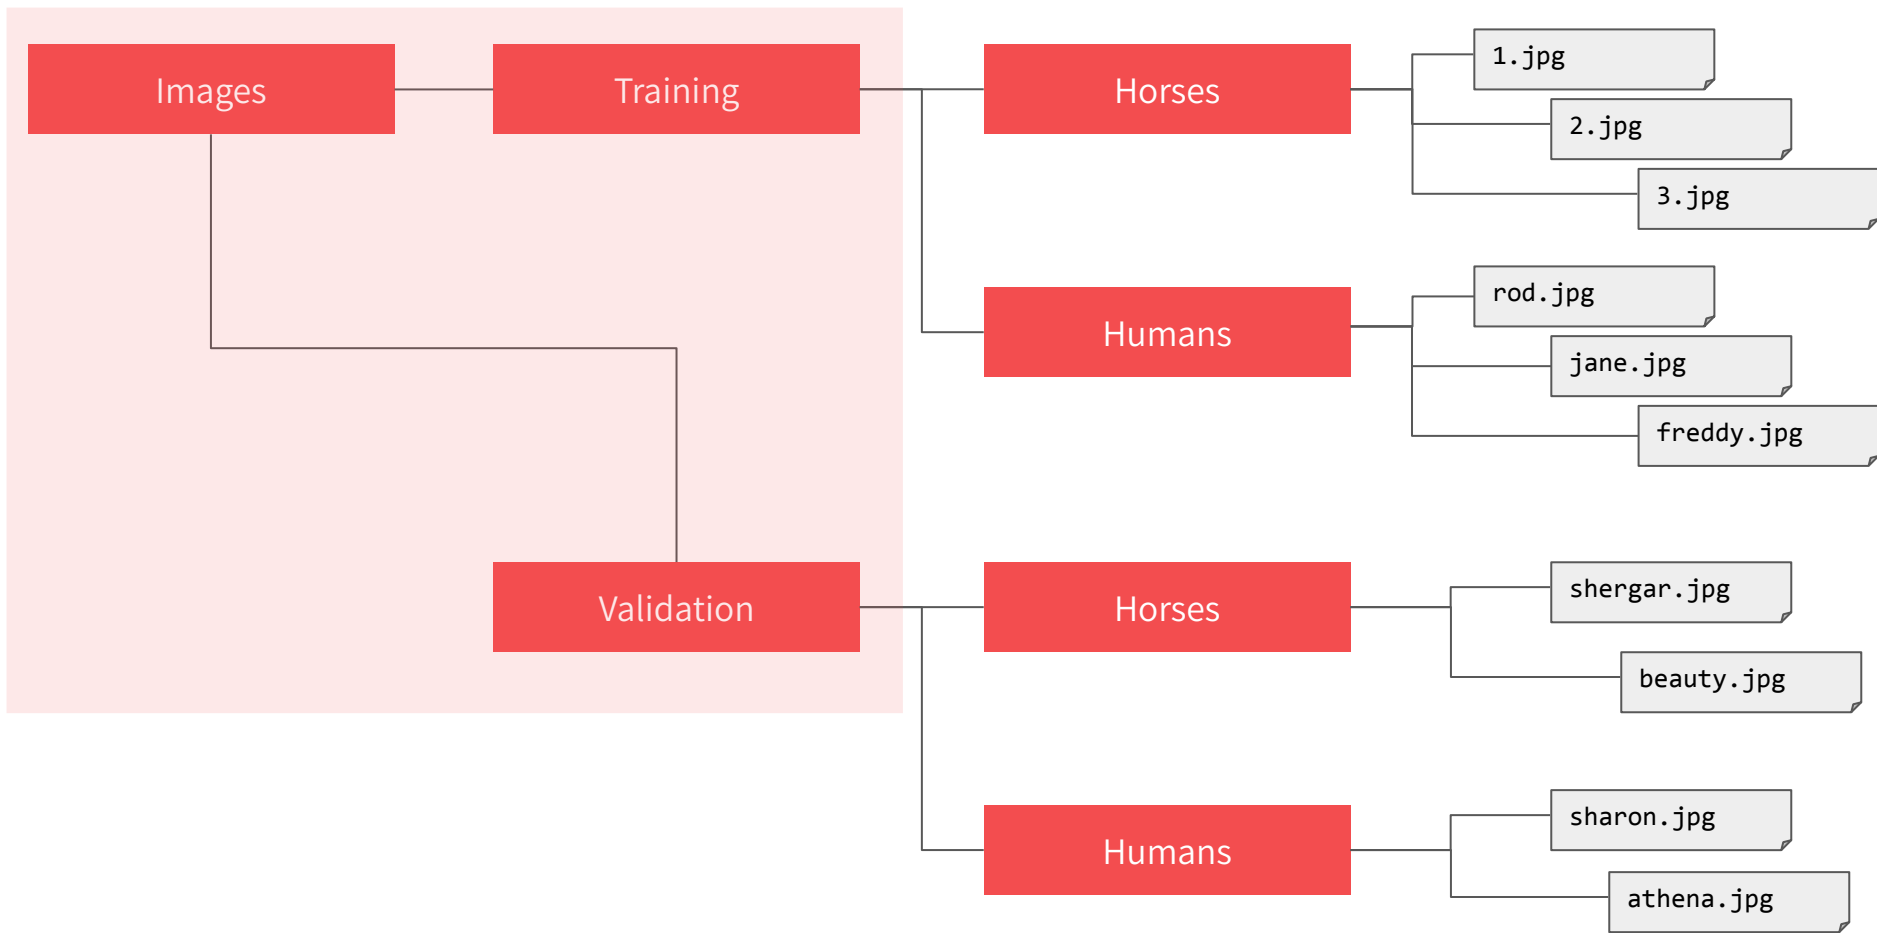
tf.data API

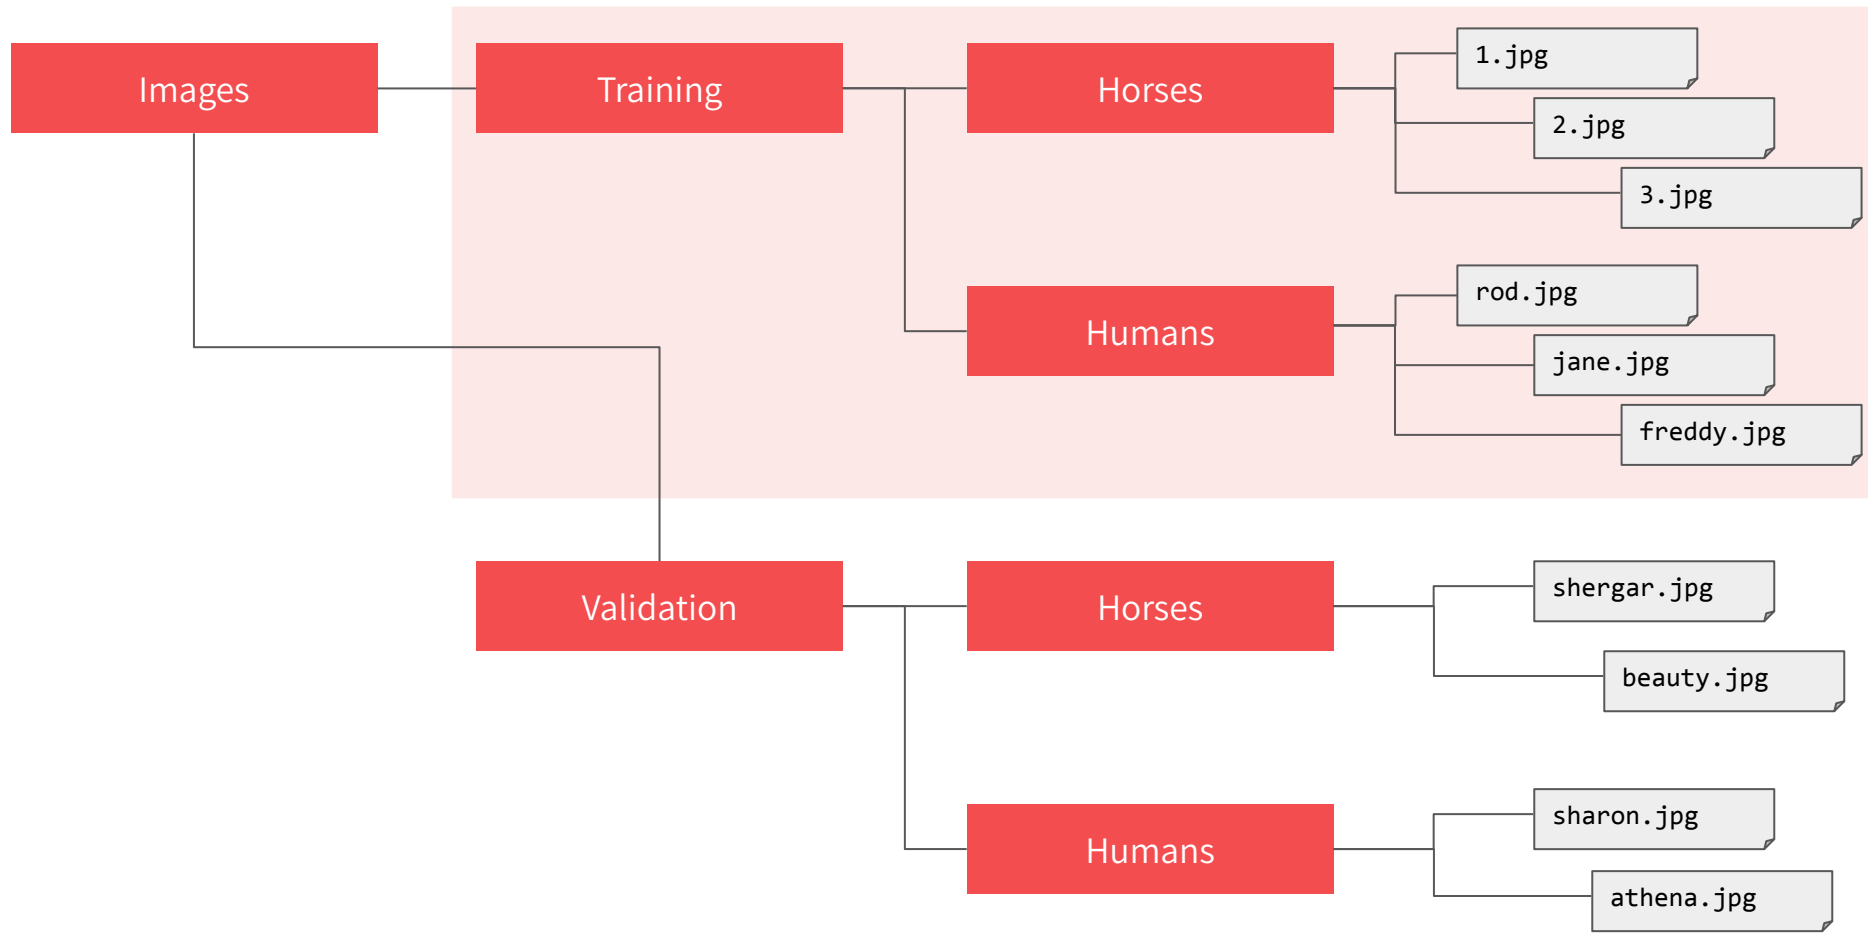



tf.keras.utils.image_dataset_from_directory











```
train_dataset = tf.keras.utils.image_dataset_from_directory(  
    TRAIN_DIR,  
    image_size=(300, 300),  
    batch_size=128,  
    label_mode='binary'  
)
```



```
train_dataset = tf.keras.utils.image_dataset_from_directory(  
    TRAIN_DIR,  
    image_size=(300, 300),  
    batch_size=128,  
    label_mode='binary'  
)
```



```
train_dataset = tf.keras.utils.image_dataset_from_directory(  
    TRAIN_DIR,  
    image_size=(300, 300),  
    batch_size=128,  
    label_mode='binary'  
)
```



```
train_dataset = tf.keras.utils.image_dataset_from_directory(  
    TRAIN_DIR,  
    image_size=(300, 300),  
    batch_size=128,  
    label_mode='binary'  
)
```



tf.data.Dataset

https://www.tensorflow.org/api_docs/python/tf/data/Dataset

```
rescale_layer = tf.keras.layers.Rescaling(scale=1./255)
```

```
train_dataset_scaled = train_dataset.map(  
    lambda image, label: (rescale_layer(image), label))
```




```
rescale_layer = tf.keras.layers.Rescaling(scale=1./255)

train_dataset_scaled = train_dataset.map(
    lambda image, label: (rescale_layer(image), label))
```



```
rescale_layer = tf.keras.layers.Rescaling(scale=1./255)

train_dataset_scaled = train_dataset.map(
    lambda image, label: (rescale_layer(image), label))
```



```
train_dataset_final = (train_dataset_scaled  
                        .cache()  
                        .shuffle(buffer_size=1000)  
                        .prefetch(buffer_size=tf.data.AUTOTUNE))
```

```
train_dataset_final = (train_dataset_scaled  
                        .cache()  
                        .shuffle(buffer_size=1000)  
                        .prefetch(buffer_size=tf.data.AUTOTUNE))
```

```
train_dataset_final = (train_dataset_scaled  
                        .cache()  
                        .shuffle(buffer_size=1000)  
                        .prefetch(buffer_size=tf.data.AUTOTUNE))
```

```
train_dataset_final = (train_dataset_scaled  
                        .cache()  
                        .shuffle(buffer_size=1000)  
                        .prefetch(buffer_size=tf.data.AUTOTUNE))
```



```
validation_dataset = tf.keras.utils.image_dataset_from_directory(
    VAL_DIR,
    image_size=(300, 300),
    batch_size=32,
    label_mode='binary'
)

validation_dataset_scaled = validation_dataset.map(lambda image, label:
(rescale_layer(image), label))

# Configure the validation dataset
validation_dataset_final = (validation_dataset_scaled
    .cache()
    .prefetch(buffer_size=tf.data.AUTOTUNE))
```



```
validation_dataset = tf.keras.utils.image_dataset_from_directory(
    VAL_DIR,
    image_size=(300, 300),
    batch_size=128,
    label_mode='binary'
)
```

```
validation_dataset_scaled = validation_dataset.map(lambda image, label:
    (rescale_layer(image), label))
```

```
# Configure the validation dataset
```

```
validation_dataset_final = (validation_dataset_scaled
    .cache()
    .prefetch(buffer_size=tf.data.AUTOTUNE))
```




```
model = tf.keras.models.Sequential([
    tf.keras.layers.Input(shape=(300, 300, 3)),
    tf.keras.layers.Conv2D(16, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Input(shape=(300, 300, 3)),
    tf.keras.layers.Conv2D(16, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Input(shape=(300, 300, 3)),
    tf.keras.layers.Conv2D(16, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Input(shape=(300, 300, 3)),
    tf.keras.layers.Conv2D(16, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

Layer (type)	Output Shape	Param #
conv2d_5 (Conv2D)	(None, 298, 298, 16)	448
max_pooling2d_5 (MaxPooling2D)	(None, 149, 149, 16)	0
conv2d_6 (Conv2D)	(None, 147, 147, 32)	4640
max_pooling2d_6 (MaxPooling2D)	(None, 73, 73, 32)	0
conv2d_7 (Conv2D)	(None, 71, 71, 64)	18496
max_pooling2d_7 (MaxPooling2D)	(None, 35, 35, 64)	0
flatten_1 (Flatten)	(None, 78400)	0
dense_2 (Dense)	(None, 512)	40141312
dense_3 (Dense)	(None, 1)	513
Total params: 40,165,409		
Trainable params: 40,165,409		
Non-trainable params: 0		

```
model.compile(loss='binary_crossentropy',  
              optimizer=tf.keras.optimizers.RMSprop(learning_rate=0.001),  
              metrics=['accuracy'])
```

<https://youtu.be/zLRB4oupj6g>



Machine Learning



0:06 / 8:58



2.1.4 Gradient Descent in Practice II Learning Rate by Andrew Ng

```
history = model.fit(  
    train_dataset_final,  
    epochs=15,  
    validation_data=validation_dataset_final,  
    verbose=2)
```



```
history = model.fit(  
    train_dataset_final,  
    epochs=15,  
    validation_data=validation_dataset_final,  
    verbose=2)
```



```
history = model.fit(  
    train_dataset_final,  
    epochs=15,  
    validation_data=validation_dataset_final,  
    verbose=2)
```



```
history = model.fit(  
    train_dataset_final,  
    epochs=15,  
    validation_data=validation_dataset_final,  
    verbose=2)
```



```
history = model.fit(  
    train_dataset_final,  
    epochs=15,  
    validation_data=validation_dataset_final,  
    verbose=2)
```

```
from google.colab import files

uploaded = files.upload()

for filename in uploaded.keys():

    # predicting images
    path = '/content/' + filename
    image = tf.keras.utils.load_img(path, target_size=(300, 300))
    image = tf.keras.utils.img_to_array(image)
    image = rescale_layer(image)
    image = np.expand_dims(image, axis=0)

    prediction = model.predict(image, verbose=0)[0][0]
    print(f'\nmodel output: {prediction}')

    if prediction > 0.5:
        print(filename + " is a human")
    else:
        print(filename + " is a horse")
```



```
from google.colab import files
```

```
uploaded = files.upload()
```

```
for filename in uploaded.keys():
```

```
    # predicting images
```

```
    path = '/content/' + filename
```

```
    image = tf.keras.utils.load_img(path, target_size=(300, 300))
```

```
    image = tf.keras.utils.img_to_array(image)
```

```
    image = rescale_layer(image)
```

```
    image = np.expand_dims(image, axis=0)
```

```
    prediction = model.predict(image, verbose=0)[0][0]
```

```
    print(f'\nmodel output: {prediction}')
```

```
    if prediction > 0.5:
```

```
        print(filename + " is a human")
```

```
    else:
```

```
        print(filename + " is a horse")
```



```
from google.colab import files
```

```
uploaded = files.upload()
```

```
for filename in uploaded.keys():
```

```
    # predicting images
```

```
    path = '/content/' + filename
```

```
    image = tf.keras.utils.load_img(path, target_size=(300, 300))
```

```
    image = tf.keras.utils.img_to_array(image)
```

```
    image = rescale_layer(image)
```

```
    image = np.expand_dims(image, axis=0)
```

```
    prediction = model.predict(image, verbose=0)[0][0]
```

```
    print(f'\nmodel output: {prediction}')
```

```
    if prediction > 0.5:
```

```
        print(filename + " is a human")
```

```
    else:
```

```
        print(filename + " is a horse")
```



```
from google.colab import files

uploaded = files.upload()

for filename in uploaded.keys():

    # predicting images
    path = '/content/' + filename
    image = tf.keras.utils.load_img(path, target_size=(300, 300))
    image = tf.keras.utils.img_to_array(image)
    image = rescale_layer(image)
    image = np.expand_dims(image, axis=0)

    prediction = model.predict(image, verbose=0)[0][0]
    print(f'\nmodel output: {prediction}')

    if prediction > 0.5:
        print(filename + " is a human")
    else:
        print(filename + " is a horse")
```

